

Research Statement

Jang-Hyun Kim, PhD candidate, Computer Science Dept. at SNU

The primary ambition of my research is to develop **sustainable AI systems that learn and adapt efficiently over long time horizons**, much like humans. I envision a future where AI systems seamlessly interact with humans and environments, forming lifelong memories and achieving meaningful, long-term engagement. However, current deep learning approaches require vast amount of labeled data to generalize on new environments, while demanding substantial memory resources during inference. This data-intensive and memory-hungry paradigm limits the applicability of AI systems to short-term engagement, raising sustainability challenge in terms of labeling efforts, computational costs, and memory consumption.

To address these limitations, my research focuses on enhancing the efficiency of neural networks during both training and inference through a data-centric approach. Specifically, my work revolves around three key themes: (1) context and dataset compression, (2) synthetic training data generation, and (3) characterization of data using relational structures. These themes are integrated into **a data optimization framework that leverages feedback from the neural network itself**, effectively harnessing the varying information densities of data while minimizing reliance on human intervention. Through this framework, I have introduced novel methods for memory-efficient inference and label-efficient training, paving the way for sustainable AI systems. In the following sections, I will highlight the key contributions of each research theme, outline my ongoing projects, and discuss future directions.

Research Progress

1 Context/Dataset Compression

Humans process real-time information selectively, forming own compressed long-term memories to enable efficient learning and lifelong adaptation to the environment. In contrast, current AI systems are inefficient in their data management. They require terabytes of storage for training data and tens of gigabytes for storing context features, as seen with Transformer models.

My primary research focuses on addressing these challenges by developing methods for compressing data and constructing long-term memory in neural networks [1, 2, 3]. I have proposed novel ideas within an integrated framework that leverages the capabilities of trained neural networks to compress information. These methods significantly enhance the efficiency of both inference and training, supporting the long-term usability of AI models in memory-constrained computing environments.

Compressed context memory for Transformers [1]. Transformer-based models, such as ChatGPT, excel at processing long-range dependencies in tokens, enabling coherent response generation based on extensive context. However, this capability comes with a significant challenge: the need to store key-value vectors for context tokens. While these vectors capture essential contextual information for response generation, their storage requirements can be prohibitive. For instance, in models like LLaMA-70B, key-value vectors for just 1,000 words demand tens of gigabytes of memory. As interactions progress, these storage demands grow rapidly, overwhelming the memory capacity of computing environments.

To address this challenge, I developed a novel online key-value compression method coupled with a compressed memory module for Transformer-based language models. My approach compresses newly introduced context into shorter key-value vectors, which are then updated in the compressed memory.


The language model accesses this compressed memory to generate contextually relevant responses. By leveraging the model’s forward pass for compression, my approach seamlessly integrates the compression process into the model’s inference pipeline. The results are promising: my method reduces the key-value cache size by 5×, while maintaining performance comparable to the original model without compression. This achievement opens up new possibilities for deploying Transformers in memory-constrained environments.

Condensing training dataset [2]. Current neural networks require vast amounts of training data, raising the key question: *What is the actual dataset size required for effective training?* If similar performance can be achieved with a smaller dataset, it would dramatically reduce the computational costs and storage required for training. To address this question, I conducted research on optimizing small synthetic datasets that match the training performance of the original dataset. Using a novel differentiable data-parameterization technique, I have successfully compressed large-scale datasets such as ImageNet with high efficiency. Notably, my approach achieves 90% of the full dataset’s training performance with only 1% of its storage size, significantly outperforming traditional data subset selection methods. Additionally, I demonstrated the effectiveness of my approach in continual learning scenarios by forming condensed example memories, further highlighting its practical utility.

Spherical data dimension reduction [3, 4]. Skeleton structures, such as human keypoints or molecules, are often represented as Cartesian products of spherical surfaces. Capturing the temporal evolution of these structures is crucial in fields like human motion analysis and molecular dynamics. However, conventional representation learning methods struggle with the inherent non-Euclidean geometry of spheres, making this task challenging. To address this issue, I developed an algorithm that optimizes a curve on spherical surfaces by minimizing the sum of geodesic distances to the data points. This approach effectively captures the underlying structure of spherical data and models the temporal evolution by projecting data onto a curve. Building on this research, I aim to further explore dataset compression methods for temporal and structured data, broadening the scope and potential impact of my work.

2 Synthetic Training Data Generation

High-quality labeled data is crucial for training high-performing AI models, but it is costly and scarce in certain domains. One promising research direction I have pursued is the use of AI models to synthesize training data. I established principles defining effective training data and developed model-driven data synthesis frameworks that create a self-reinforcing loop, enhancing overall training performance. This approach holds significant potential to drive sustained improvements in AI systems while reducing human labeling efforts.



Compressed memory system for memory-efficient Transformer inference [1]. Colored boxes represent the attention key-value (KV) pairs.


Class: Lorikeet



Class: Bottle cap



Condensed ImageNet samples [2]. My approach constructs compressed form of training data information, enabling efficient training.



Spherical principal curve [3] encodes the temporal evolution of spherical data through dimension reduction.

Saliency-guided data augmentation [5]. Conventional data augmentation methods often rely on predefined functions that randomly alter data without accounting for its individual characteristics, such as object locations in images. This results in the generation of data that are inconsistent with the assigned labels, providing false supervisory signals. To overcome this challenge, I developed a novel approach that leverages the data saliency maps obtained by the model under training. Using this saliency information, I designed an augmentation technique that mixes a data pair while preserving the most informative regions. This feedback-driven method not only improves the model’s generalization performance but also guides the model more robust to input noise. The impact of my work extends beyond image processing, with successful applications in various domains including natural language, graph, and point cloud data.

3 Characterization of Data using Relational Structure

To tackle these critical challenges in real-world applications, I developed a framework to systematically identify complex data issues. Moving beyond conventional single-feature metrics like prediction error, I proposed a method that analyzes the relationships between data points. Utilizing pretrained models, I introduced *Neural Relation Graph* [7], a fully-connected graph of data points, where each edge weight quantifies the degree of complementarity or conflict between data points. This graph provides rich information of individual data points, which I leveraged to develop an algorithm that effectively identifies label errors as well as outliers. Additionally, I designed a visualization tool that enables users to intuitively explore these relationships in a 2D plot, thereby enhancing interpretability and usability of my method.

Puzzle Mix [5] generates synthetic training data by mixing salient regions from data pairs. I extract the saliency map from the model under training, creating a self-reinforcing loop in the training process.

Co-Mixup [6] generates a batch of synthetic data, jointly optimizing the diversity and saliency of outputs.

Neural relation graph [7] encodes relational structures within data using neural networks. By characterizing each data point within the graph, my algorithm effectively identifies outliers and mislabeled data.

Ongoing and Future Direction

In summary, my research has focused on optimizing data usage in AI systems to develop sustainable systems capable of (1) efficient memory management during inference and (2) effective learning from limited human-labeled data. I am eager to deepen and broaden my research to achieve tangible impacts in real-world applications and sciences through the following research topics:

- *What is the most effective memory structure for long-context processing?* I am interested in enabling neural networks to store information with minimal capacity requirements. This problem involves identifying the minimum description length of information within the neural network’s memory space (*i.e.*, model weights or data features). The objective is to determine the most compact memory structure, allowing neural networks to efficiently store vast volumes of information and support long-term interactions.
- Another major ambition of mine is to create *multi-modal memory systems*, drawing on my research experience across various domains, including image [2, 5, 6, 7], speech [2, 10], natural language [1, 7], and multi-joint dynamics [3, 9]. A particular focus is video data which presents substantial challenges due to its extensive size and information sparsity, requiring innovative compression techniques. The goal is to expand my research to develop AI systems capable of long-term video interactions, thereby enhancing the AI’s utility in real-world applications.
- I am also interested in enabling AI to *identify causal relationships among data*. In my recent work on *Targeted Cause Discovery* [8], I explored a data-driven learning approach for causal discovery, successfully identifying novel causal relationships among human genes. Moving forward, I aim to improve AI models’ capability to accurately identify causal relationships in complex real-world environments (*e.g.*, egocentric videos and time-series forecasting), thereby enhancing their reasoning and memorization capabilities.

References

- [1] **Jang-Hyun Kim**, Junyoung Yeom, Sangdoo Yun, Hyun Oh Song. “Compressed Context Memory for Online Language Model Interaction”. In *ICLR*. 2024.
- [2] **Jang-Hyun Kim**, Jinuk Kim, Seong Joon Oh, Sangdoo Yun, Hwanjun Song, Joonhyun Jeong, Jung-Woo Ha, Hyun Oh Song. “Dataset Condensation via Efficient Synthetic-Data Parameterization”. In *ICML*. 2022.
- [3] Jongmin Lee*, **Jang-Hyun Kim***, Hee-Seok Oh (*: equal contribution). “Spherical Principal Curves”. In *TPAMI*. 2021.
- [4] Jongmin Lee, **Jang-Hyun Kim**, Hee-Seok Oh. “spherepc: An R Package for Dimension Reduction on a Sphere”. In *R Journal*. 2022.
- [5] **Jang-Hyun Kim**, Wonho Choo, Hosan Jeong, Hyun Oh Song. “Puzzle Mix: Exploiting Saliency and Local Statistics for Optimal Mixup”. In *ICML*. 2020.
- [6] **Jang-Hyun Kim**, Wonho Choo, Hosan Jeong, Hyun Oh Song. “Co-Mixup: Saliency Guided Joint Mixup with Supermodular Diversity”. In *ICLR (oral presentation)*. 2021.
- [7] **Jang-Hyun Kim**, Sangdoo Yun, Hyun Oh Song. “Neural Relation Graph: A Unified Framework for Identifying Label Noise and Outlier Data”. In *NeurIPS*. 2023.
- [8] **Jang-Hyun Kim**, Claudia Skok Gibbs, Sangdoo Yun, Hyun Oh Song, Kyunghyun Cho. “Targeted Cause Discovery with Data-Driven Learning”. *arXiv preprint (under review)*. 2024.
- [9] Gaon An*, Seungyong Moon*, **Jang-Hyun Kim**, Hyun Oh Song (*: equal contribution). “Uncertainty-Based Offline Reinforcement Learning with Diversified Q-ensemble”. In *NeurIPS*. 2021.
- [10] **Jang-Hyun Kim***, Jaejun Yoo*, Sanghyuk Chun, Adrian Kim, Jung-Woo Ha (*: equal contribution). “Multi-Domain Processing via Hybrid Denoising Networks for Speech Enhancement”. *arXiv preprint*. 2018.